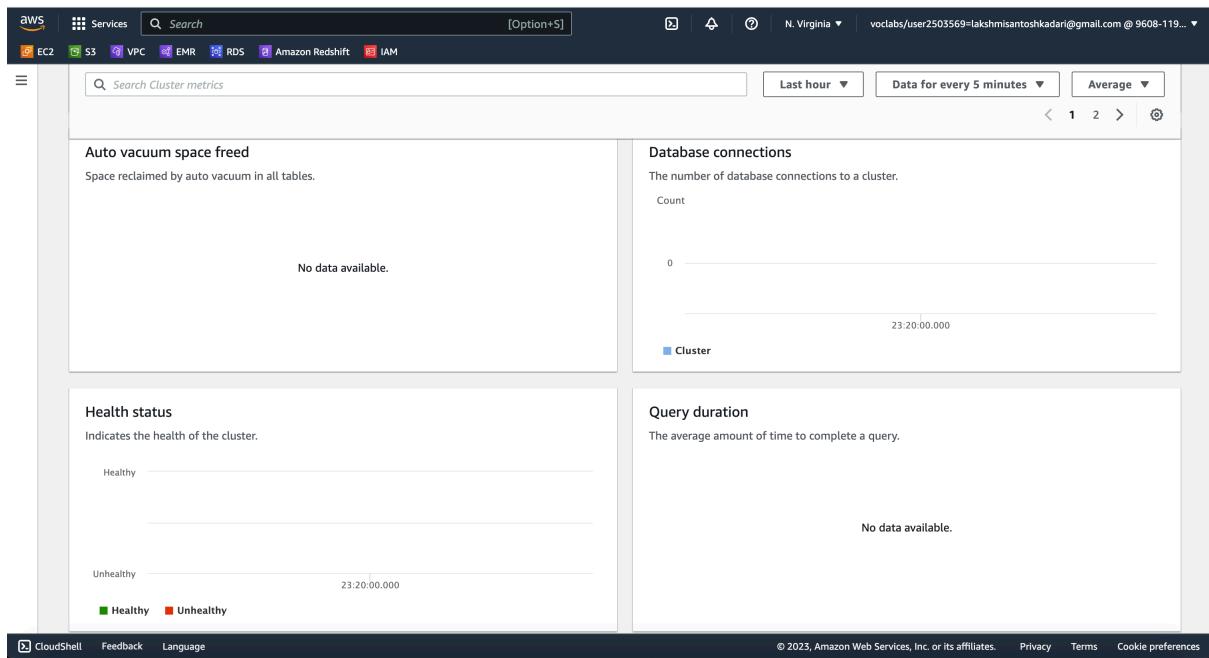


RedShift Setup

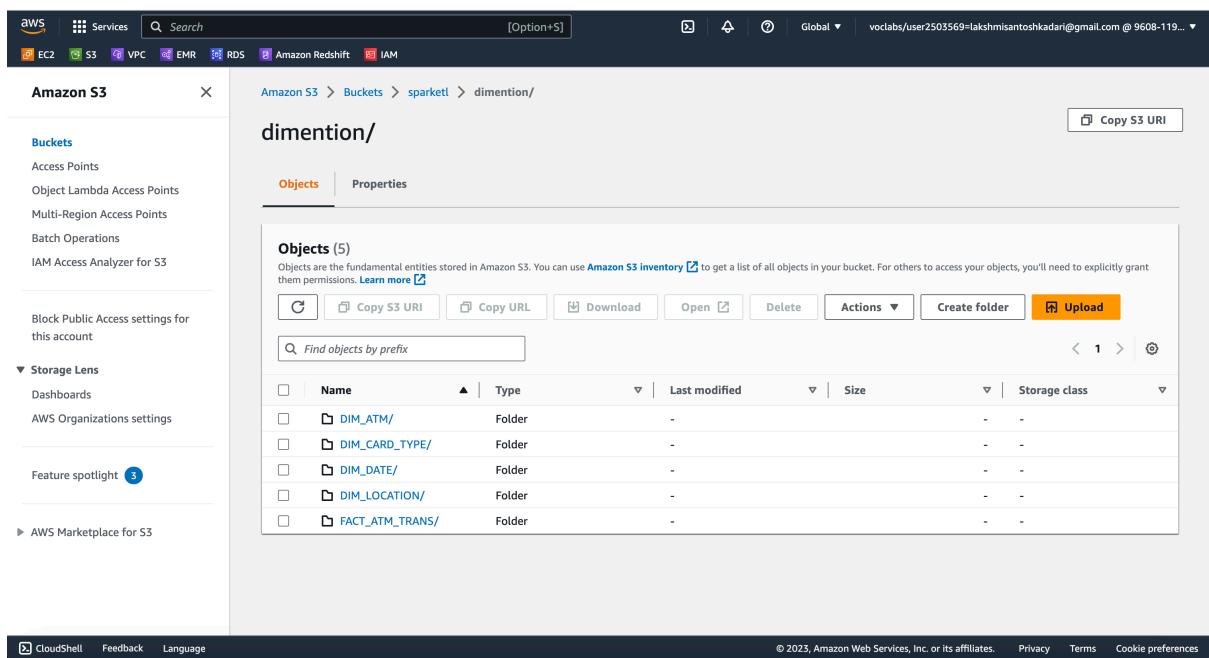
1. Amazon Redshift creation

The screenshot shows the AWS Management Console with the URL <https://console.aws.amazon.com/redshift/clusters?region=N.%20Virginia>. The cluster name is etlsparkproject. The cluster identifier is etlsparkproject. The status is Available. The node type is dc2.large. The endpoint is etlsparkproject.c2dtrd1wc9rq.us-east-1.redshift.amazonaws.com:5439/dev. The JDBC URL is jdbc:redshift://etlsparkproject.c2dtrd1wc9rq.us-east-1.redshift.amazonaws.com:5439/dev. The ODBC URL is Driver=(Amazon Redshift (x64));Server=etlsparkproject.c2dtrd1wc9rq.us-east-1.redshift.amazonaws.com;Database=dev. The cluster was created on July 22, 2023, at 23:23 (UTC+05:30). It has 2 nodes and is in a Multi-AZ configuration. The configuration is Production.

The screenshot shows the AWS Management Console with the URL <https://console.aws.amazon.com/redshift/metrics?clusterName=etlsparkproject®ion=N.%20Virginia>. The cluster metrics page displays various metrics for the etlsparkproject cluster. It includes sections for Alarms (0), Events (1), Cluster metrics (11/11), CPU utilization, Percentage disk space used, Auto vacuum space freed, and Database connections. The CPU utilization chart shows the percentage of CPU utilization over time, with data points for Leader, Compute-0, and Compute-1 nodes. The Percentage disk space used chart indicates that no data is available. The Auto vacuum space freed chart shows the space reclaimed by auto vacuum in all tables. The Database connections chart shows the number of database connections to a cluster.

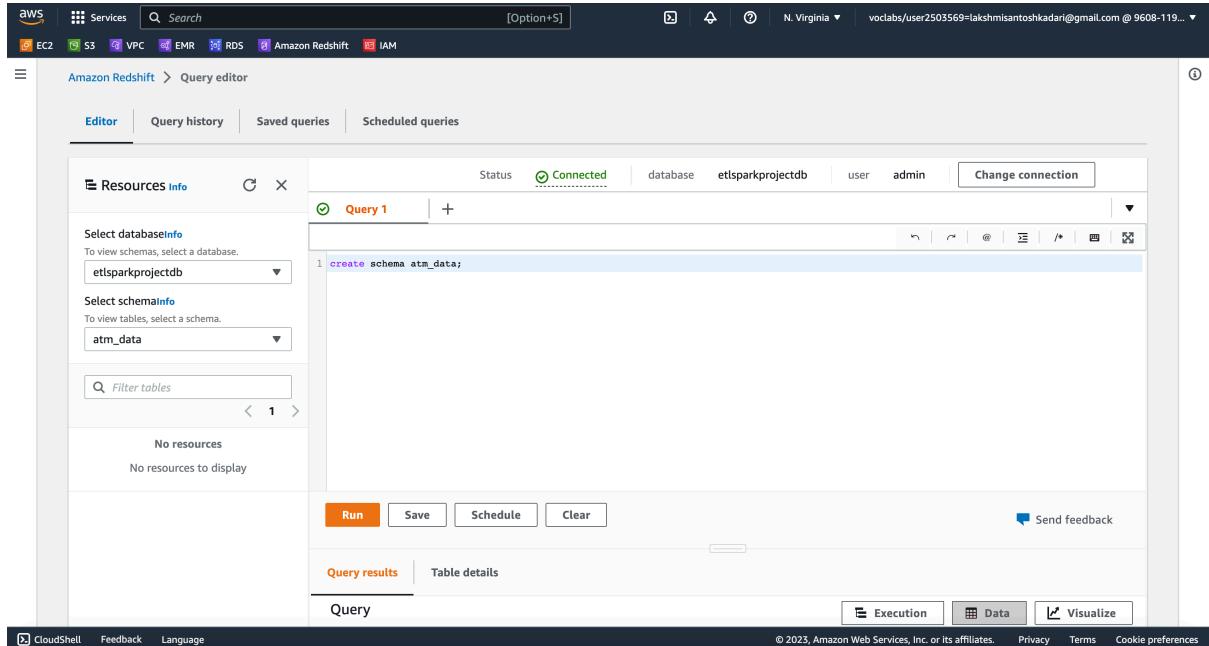


Viewing all the dimension and fact data in Amazon S3:



Query to create a schema for the dimension and fact tables:

create schema atm_data;



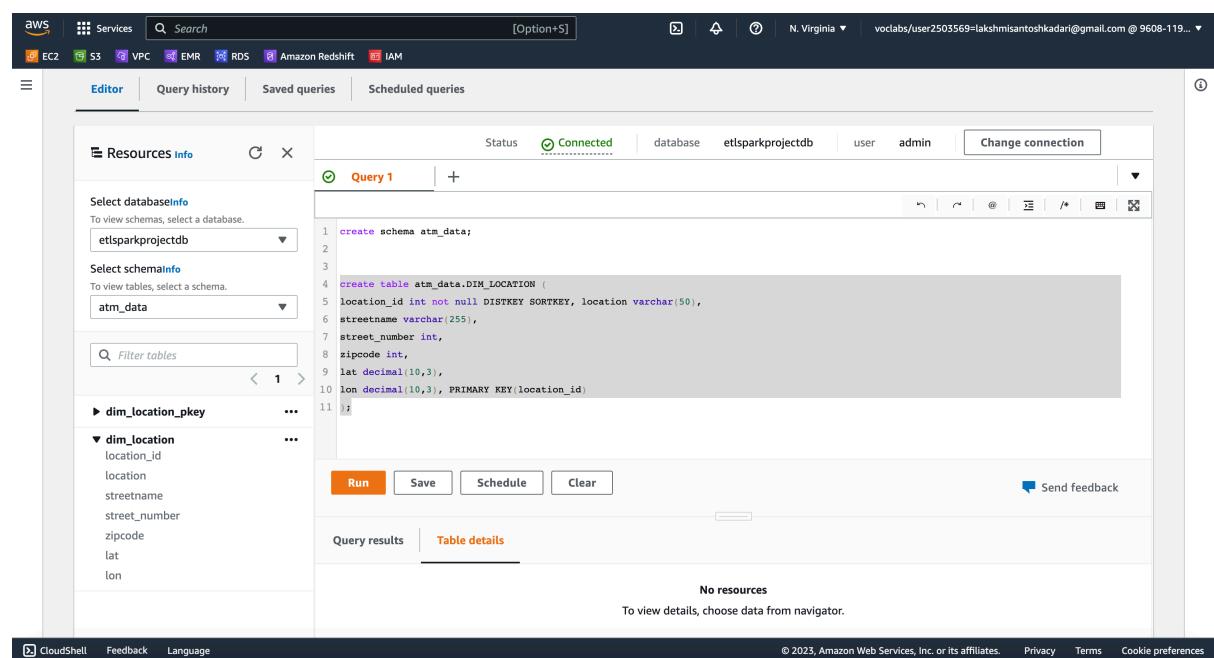
The screenshot shows the AWS Management Console interface for Amazon Redshift. At the top, the navigation bar includes the AWS logo, Services (with EC2, S3, VPC, EMR, RDS, Amazon Redshift, IAM), a search bar, and account information (vocabs/user2503569=lakshmisantoshkadari@gmail.com @ 9608-119...). Below the navigation bar, the title "Amazon Redshift > Query editor" is displayed. The main area is the "Query editor". On the left, there is a sidebar titled "Resources Info" with sections for "Select databaseInfo" (set to etlsparkprojectdb) and "Select schemaInfo" (set to atm_data). The main query editor window shows a single query: "create schema atm_data;". The status bar at the bottom indicates "Connected" to the database etlsparkprojectdb, user admin, and shows "No resources" found. The footer contains links for CloudShell, Feedback, Language, and various AWS terms like Privacy, Terms, and Cookie preferences.

```
create schema atm_data;
```

Queries to create the various dimension and fact tables with appropriate primary and foreign keys:

Creating location dimension table

```
create table atm_data.DIM_LOCATION (
    location_id int not null DISTKEY SORTKEY, location varchar(50),
    streetname varchar(255),
    street_number int,
    zipcode int,
    lat decimal(10,3),
    lon decimal(10,3), PRIMARY KEY(location_id)
);
```



The screenshot shows the AWS Redshift Query Editor interface. The top navigation bar includes 'Services' (with 'Search' and 'Option+S' dropdown), 'EC2', 'S3', 'VPC', 'EMR', 'RDS', 'Amazon Redshift', and 'IAM'. The main area has tabs for 'Editor' (selected), 'Query history', 'Saved queries', and 'Scheduled queries'. On the left, there's a sidebar titled 'Resources' with sections for 'Info', 'Select databaseinfo' (set to 'etlsparkprojectdb'), 'Select schemainfo' (set to 'atm_data'), and a 'Filter tables' search bar. Below these are two expandable sections: 'dim_location_pkey' and 'dim_location', which lists columns: location_id, location, streetname, street_number, zipcode, lat, and lon. The central workspace contains a query editor with the SQL code provided above. The status bar at the bottom shows 'CloudShell', 'Feedback', 'Language', and copyright information: '© 2023, Amazon Web Services, Inc. or its affiliates.' followed by 'Privacy', 'Terms', and 'Cookie preferences'.

```
1 create schema atm_data;
2
3
4 create table atm_data.DIM_LOCATION (
5     location_id int not null DISTKEY SORTKEY, location varchar(50),
6     streetname varchar(255),
7     street_number int,
8     zipcode int,
9     lat decimal(10,3),
10    lon decimal(10,3), PRIMARY KEY(location_id)
11 );
```

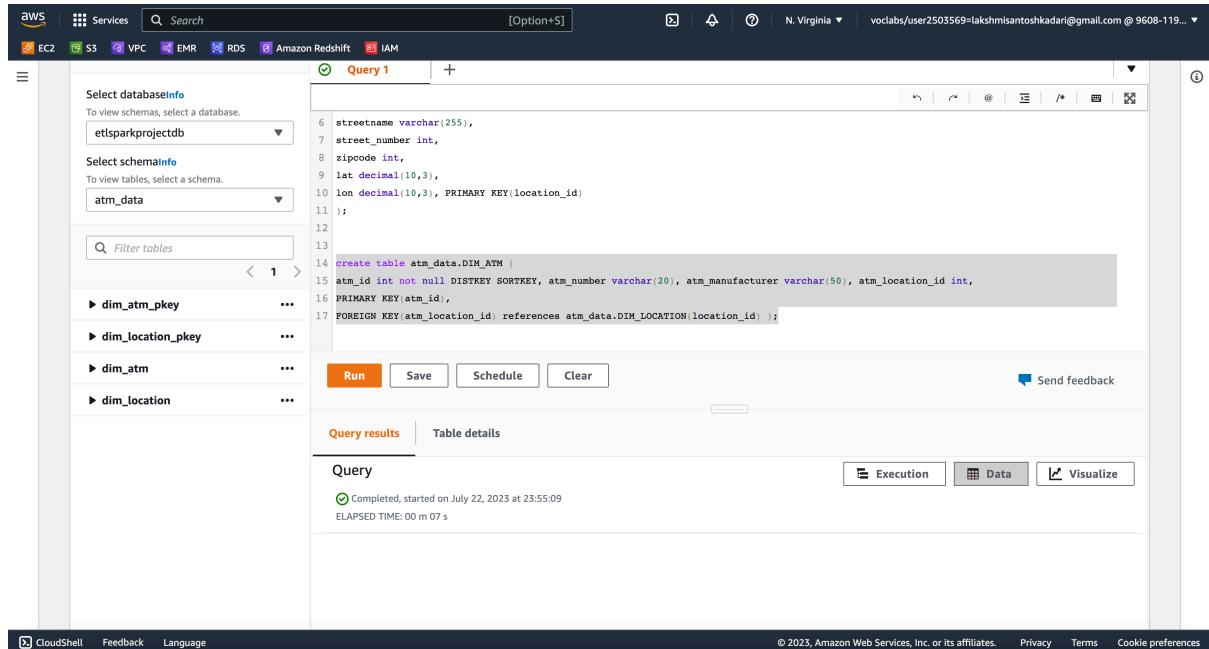
Creating atm dimension table

```
create table atm_data.DIM_ATM (
```

```
atm_id int not null DISTKEY SORTKEY, atm_number varchar(20), atm_manufacturer  
varchar(50), atm_location_id int,
```

```
PRIMARY KEY(atm_id),
```

```
FOREIGN KEY(atm_location_id) references atm_data.DIM_LOCATION(location_id);
```



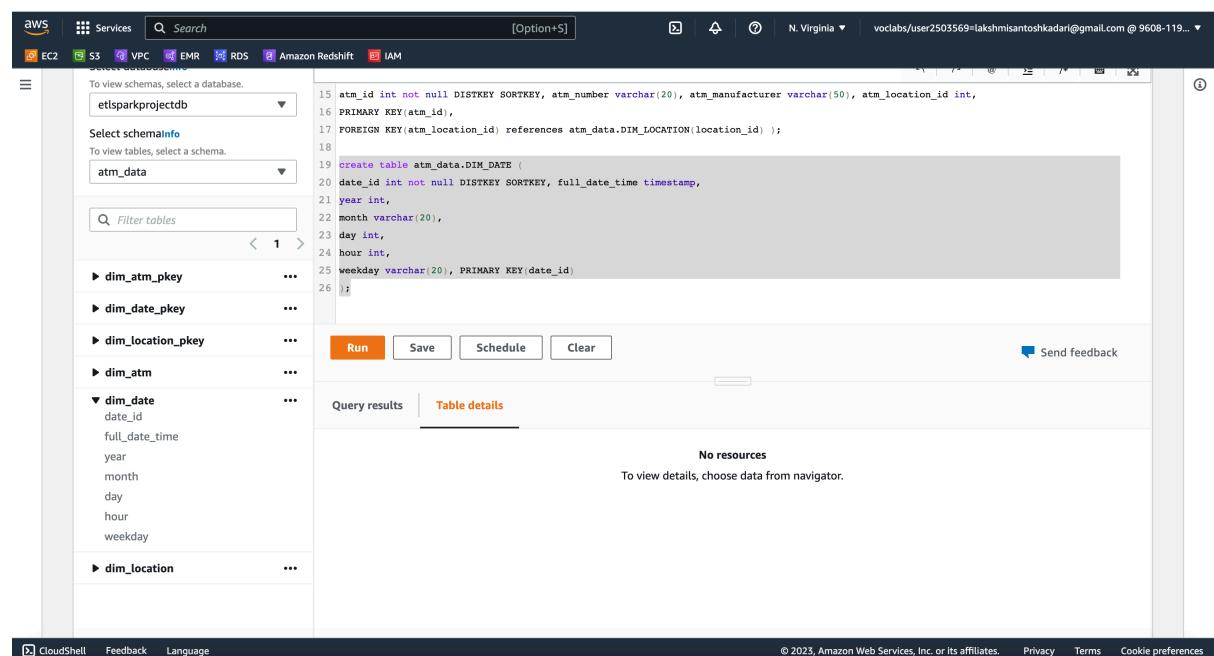
The screenshot shows the AWS CloudShell interface. On the left, there's a sidebar with database and schema selection dropdowns: 'Select databaseinfo' set to 'etlsparkprojectdb' and 'Select schemainfo' set to 'atm_data'. Below these are buttons for 'Filter tables' and a list of tables: 'dim_atm_pkey', 'dim_location_pkey', 'dim_atm', and 'dim_location'. The main area contains a query editor titled 'Query 1'. The code entered is:

```
6 streetname varchar(255),  
7 street_number int,  
8 zipcode int,  
9 lat decimal(10,3),  
10 lon decimal(10,3), PRIMARY KEY(location_id),  
11 ;  
12  
13  
14 create table atm_data.DIM_ATM :  
15 atm_id int not null DISTKEY SORTKEY, atm_number varchar(20), atm_manufacturer varchar(50), atm_location_id int,  
16 PRIMARY KEY(atm_id),  
17 FOREIGN KEY(atm_location_id) references atm_data.DIM_LOCATION(location_id) ;
```

Below the code are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. To the right of the code area is a 'Send feedback' button. At the bottom of the editor, there are tabs for 'Query results' and 'Table details', and buttons for 'Execution', 'Data', and 'Visualize'. The status bar at the bottom indicates the query completed successfully on July 22, 2023, at 23:55:09, with an elapsed time of 00 m 07 s.

Creating date dimension table

```
create table atm_data.DIM_DATE (  
  
date_id int not null DISTKEY SORTKEY, full_date_time timestamp,  
year int,  
month varchar(20),  
  
day int,  
hour int,  
weekday varchar(20), PRIMARY KEY(date_id)  
  
);
```

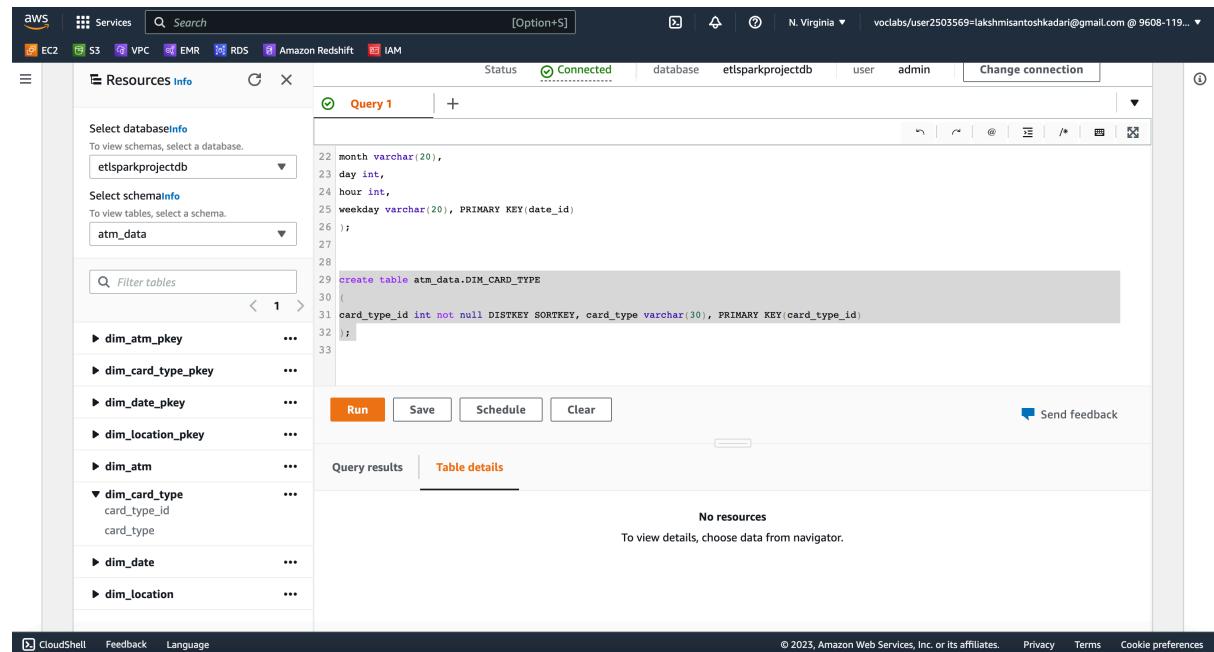


The screenshot shows the AWS CloudShell interface. The top navigation bar includes 'Services' (with 'Search' and '[Option+S]'), 'EC2', 'S3', 'VPC', 'EMR', 'RDS', 'Amazon Redshift', and 'IAM'. The top right shows 'N. Virginia' and a user session. The main area has a dark background with white text. On the left, there's a sidebar titled 'Select database' containing 'etlsparkprojectdb' and a dropdown for 'atm_data'. Below that is a 'Filter tables' search bar. The main workspace contains the SQL code for creating the 'DIM_DATE' table. At the bottom are 'Run', 'Save', 'Schedule', and 'Clear' buttons, along with a 'Send feedback' link. The status bar at the bottom includes links for 'CloudShell', 'Feedback', 'Language', and copyright information: '© 2023, Amazon Web Services, Inc. or its affiliates.' followed by 'Privacy', 'Terms', and 'Cookie preferences'.

```
15 atm_id int not null DISTKEY SORTKEY, atm_number varchar(20), atm_manufacturer varchar(50), atm_location_id int,  
16 PRIMARY KEY(atm_id),  
17 FOREIGN KEY(atm_location_id) references atm_data.DIM_LOCATION(location_id) );  
18  
19 create table atm_data.DIM_DATE (  
20 date_id int not null DISTKEY SORTKEY, full_date_time timestamp,  
21 year int,  
22 month varchar(20),  
23 day int,  
24 hour int,  
25 weekday varchar(20), PRIMARY KEY(date_id)  
26 );
```

Creating card type dimension table

```
create table atm_data.DIM_CARD_TYPE (
    card_type_id int not null DISTKEY SORTKEY, card_type varchar(30),
    PRIMARY KEY(card_type_id)
);
```



The screenshot shows the AWS CloudShell interface. The top navigation bar includes 'Search' and 'Option+S'. Below it, the status bar shows 'Connected' to 'etlsparkprojectdb' database, 'user admin', and the connection string 'voclabs/user2503569=lakshmisantoshkadari@gmail.com @ 9608-119...'. The main area has a sidebar with 'Resources info' and dropdowns for 'Select databaseInfo' (set to 'etlsparkprojectdb') and 'Select schemaInfo' (set to 'atm_data'). A central query editor window titled 'Query 1' contains the SQL code for creating the 'DIM_CARD_TYPE' table. The code is as follows:

```
22 month varchar(20),
23 day int,
24 hour int,
25 weekday varchar(20), PRIMARY KEY(date_id)
26 );
27
28
29 create table atm_data.DIM_CARD_TYPE
30 (
31 card_type_id int not null DISTKEY SORTKEY, card_type varchar(30), PRIMARY KEY(card_type_id)
32 );
33
```

Below the query editor are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. To the right is a 'Send feedback' link. At the bottom of the editor are tabs for 'Query results' and 'Table details', with 'Table details' currently selected. The message 'No resources' is displayed under the table details tab. The footer of the interface includes links for 'CloudShell', 'Feedback', 'Language', and copyright information: '© 2023, Amazon Web Services, Inc. or its affiliates.' followed by 'Privacy', 'Terms', and 'Cookie preferences'.

Creating atm transactions fact table

```
create table atm_data.FACT_ATM_TRANS (
    trans_id bigint not null DISTKEY SORTKEY,
    atm_id int,
    weather_loc_id int,
    date_id int,
    card_type_id int,
    atm_status varchar(20),
    currency varchar(10),
    service varchar(20),
    transaction_amount int,
    message_code varchar(225),
    message_text varchar(225),
    rain_3h decimal(10,3),
    clouds_all int,
    weather_id int,
    weather_main varchar(50),
    weather_description varchar(255),
    PRIMARY KEY(trans_id),
    FOREIGN KEY(weather_loc_id) references atm_data.DIM_LOCATION(location_id),
    FOREIGN KEY(atm_id) references atm_data.DIM_ATM(atm_id),
    FOREIGN KEY(date_id) references atm_data.DIM_DATE(date_id),
    FOREIGN KEY(card_type_id) references atm_data.DIM_CARD_TYPE(card_type_id)
);
```

The screenshot shows the AWS Redshift Query Editor interface. On the left, there's a sidebar with dropdown menus for 'Select databaseinfo' (set to 'etlsparkprojectdb') and 'Select schema/info' (set to 'atm_data'). Below these are lists of tables like 'dim_atm_pkey', 'dim_card_type_pkey', etc., and a detailed view of the 'fact_atm_trans_pkey' table. The main panel displays the SQL code for creating the 'FACT_ATM_TRANS' table:

```

35 create table atm_data.FACT_ATM_TRANS (
36     trans_id bigint not null DISTKEY SORTKEY,
37     atm_id int,
38     weather_loc_id int,
39     date_id int,
40     card_type_id int,
41     atm_status varchar(20),
42     currency varchar(10),
43     service varchar(20),
44     transaction_amount int,
45     message_code varchar(225),
46     message_text varchar(225),
47     rain_3h decimal(10,3),
48     ...

```

Below the code are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The status bar at the bottom right shows 'Send feedback'.

Loading data into a RedShift cluster from Amazon S3 bucket

Queries to copy the data from S3 bucket to the RedShift cluster into appropriate tables:

```

copy atm_data.dim_location from 's3://sparketl/dimention/DIM_LOCATION/part-00000-
ccb498ec-94dc-47a9-8761-efa53aa9f106-c000.csv'
iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
delimiter ',' region 'us-east-1'
CSV;

```

The screenshot shows the AWS Redshift Query Editor interface. The sidebar is identical to the previous one. The main panel shows a query editor with a single query named 'Query 1' containing the COPY command. The status bar at the bottom right indicates the query was completed successfully on July 23, 2023, at 00:20:33, with an elapsed time of 00 m 57 s.

```

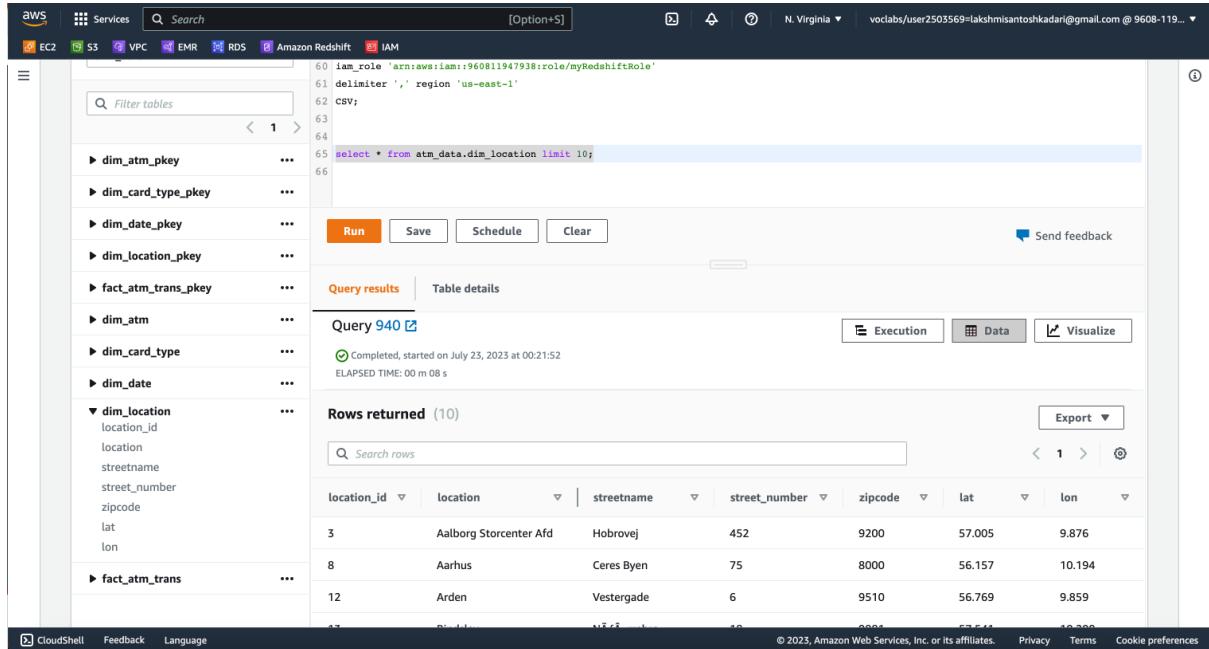
52 PRIMARY KEY(trans_id),
53 FOREIGN KEY(weather_loc_id) references atm_data.DIM_LOCATION(location_id), FOREIGN KEY(atm_id) references atm_data.DIM_ATM(atm_id),
54 FOREIGN KEY(date_id) references atm_data.DIM_DATE(date_id),
55 FOREIGN KEY(card_type_id) references atm_data.DIM_CARD_TYPE(card_type_id)
56 );
57
58
59 copy atm_data.dim_location from 's3://sparketl/dimention/DIM_LOCATION/part-00000-ccb498ec-94dc-47a9-8761-efa53aa9f106-c000.csv'
60 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
61 delimiter ',' region 'us-east-1'
62 CSV;
63
64

```

The status bar at the bottom right shows 'Send feedback'.

Verified the data to ensure data got uploaded using the below query.

select * from atm_data.dim_location limit 10;



```
60 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
61 delimiter ',' region 'us-east-1'
62 CSV;
63
64
65 select * from atm_data.dim_location limit 10;
66
```

Run Save Schedule Clear Send feedback

Query results | Table details

Query 940 []

Completed, started on July 23, 2023 at 00:21:52
ELAPSED TIME: 00 m 08 s

Rows returned (10)

location_id	location	streetname	street_number	zipcode	lat	lon
3	Aalborg Storcenter Afd	Hobrovej	452	9200	57.005	9.876
8	Aarhus	Ceres Byen	75	8000	56.157	10.194
12	Arden	Vestergade	6	9510	56.769	9.859
17	Piedmont	W. 2nd	10	8001	57.511	10.200

Export < 1 > @

CloudShell Feedback Language © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Copying the data to dim_date table

```
copy atm_data.dim_date from 's3://sparketl/dimention/DIM_DATE/part-00000-2704fc3a-d03d-4a54-90f6-bbf429f7203b-c000.csv'
iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
delimiter ',' region 'us-east-1'
timeformat 'YYYY-MM-DDTHH:MI:SS'
CSV;
```

```

62 CSV;
63
64 select * from atm_data.dim_location limit 10;
65
66
67 copy atm_data.dim_date from 's3://sparktel/dimention/DIM_DATE/part-00000-2704fc3a-d03d-4a54-90f6-bbf429f7203b-c000.csv'
68 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
69 delimiter ',' region 'us-east-1'
70 timestampformat 'YYYY-MM-DDTHH:MI:SS'
71 CSV;
72
73 select * from atm_data.dim_date limit 10;

```

Run **Save** **Schedule** **Clear**

Query results **Table details**

No resources

To view details, choose data from navigator.

Verified the data to ensure data got uploaded using the below query.

`select * from atm_data.dim_date limit 10;`

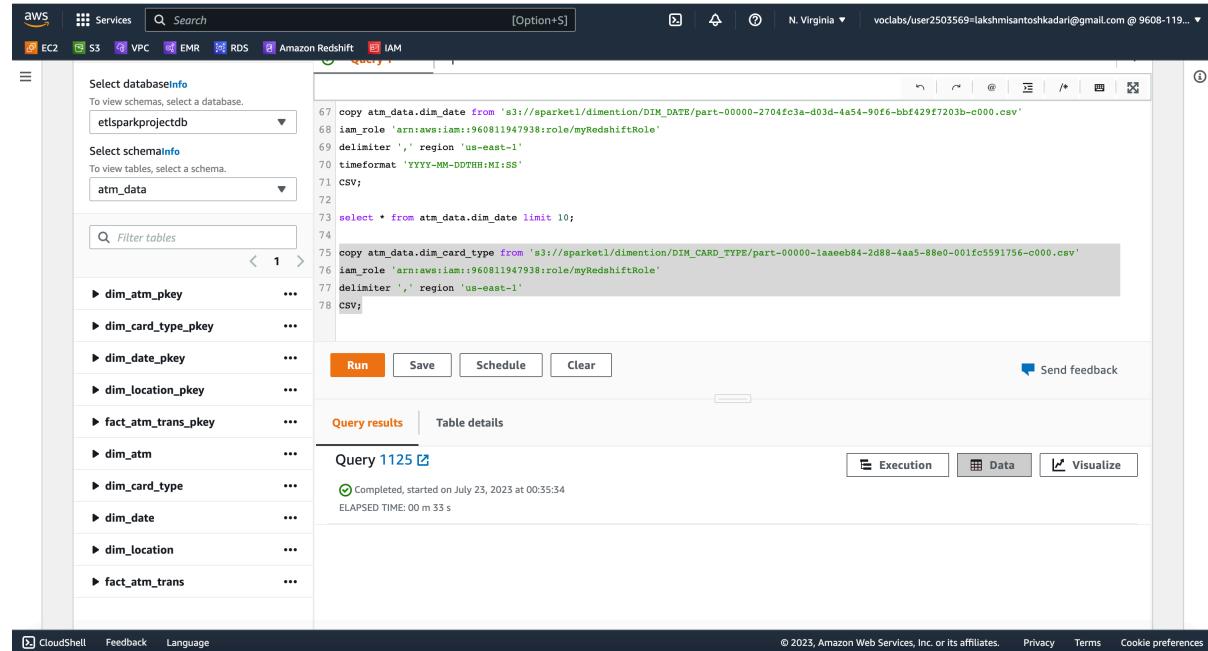
Query 1078 Completed, started on July 23, 2023 at 00:31:05 ELAPSED TIME: 00 m 08 s

Rows returned (10)

date_id	full_date_time	year	month	day	hour	weekday
1	2017-01-01 00:00:00	2017	January	1	0	Sunday
6	2017-01-01 05:00:00	2017	January	1	5	Sunday
11	2017-01-01 10:00:00	2017	January	1	10	Sunday

Copying the data to dim_card_type table

```
copy atm_data.dim_card_type from 's3://sparketl/dimention/DIM_CARD_TYPE/part-00000-1aae...  
iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
delimiter ',' region 'us-east-1'  
CSV;
```



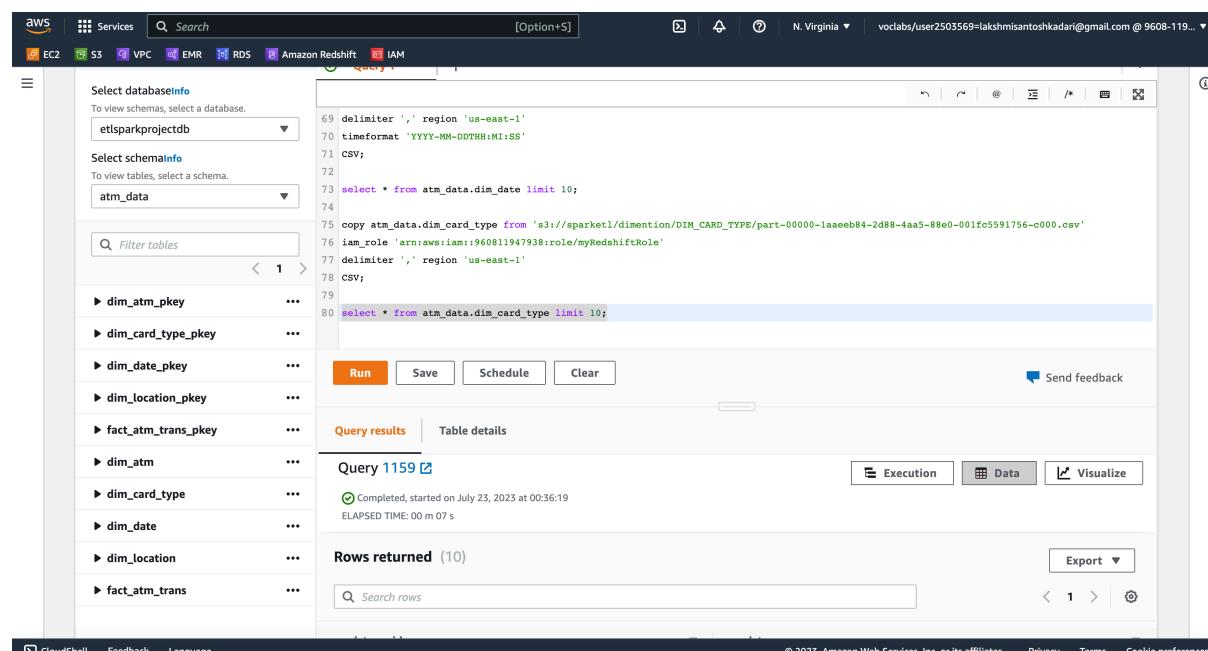
The screenshot shows the AWS Redshift Query Editor interface. On the left, there's a sidebar with dropdown menus for 'Select databaseinfo' (set to etisparkprojectdb) and 'Select schemainfo' (set to atm_data). Below these are lists of tables: dim_atm_pkey, dim_card_type_pkey, dim_date_pkey, dim_location_pkey, fact_atm_trans_pkey, dim_atm, dim_card_type, dim_date, dim_location, and fact_atm_trans. The main area contains the following SQL code:

```
67 copy atm_data.dim_date from 's3://sparketl/dimention/DIM_DATE/part-00000-2704fc3a-d03d-4a54-90f6-bbf429ff7203b-c000.csv'  
68 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
69 delimiter ',' region 'us-east-1'  
70 timeformat 'YYYY-MM-DDTHH:MI:SS'  
71 CSV;  
72  
73 select * from atm_data.dim_date limit 10;  
74  
75 copy atm_data.dim_card_type from 's3://sparketl/dimention/DIM_CARD_TYPE/part-00000-1aae...  
76 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
77 delimiter ',' region 'us-east-1'  
78 CSV;
```

Below the code, there are 'Run', 'Save', 'Schedule', and 'Clear' buttons. To the right, there are tabs for 'Query results' (selected), 'Table details', 'Execution', 'Data', and 'Visualize'. The status bar at the bottom indicates 'Completed, started on July 23, 2023 at 00:35:34' and 'ELAPSED TIME: 00 m 33 s'.

Verified the data to ensure data got uploaded using the below query.

```
select * from atm_data.dim_card_type limit 10;
```



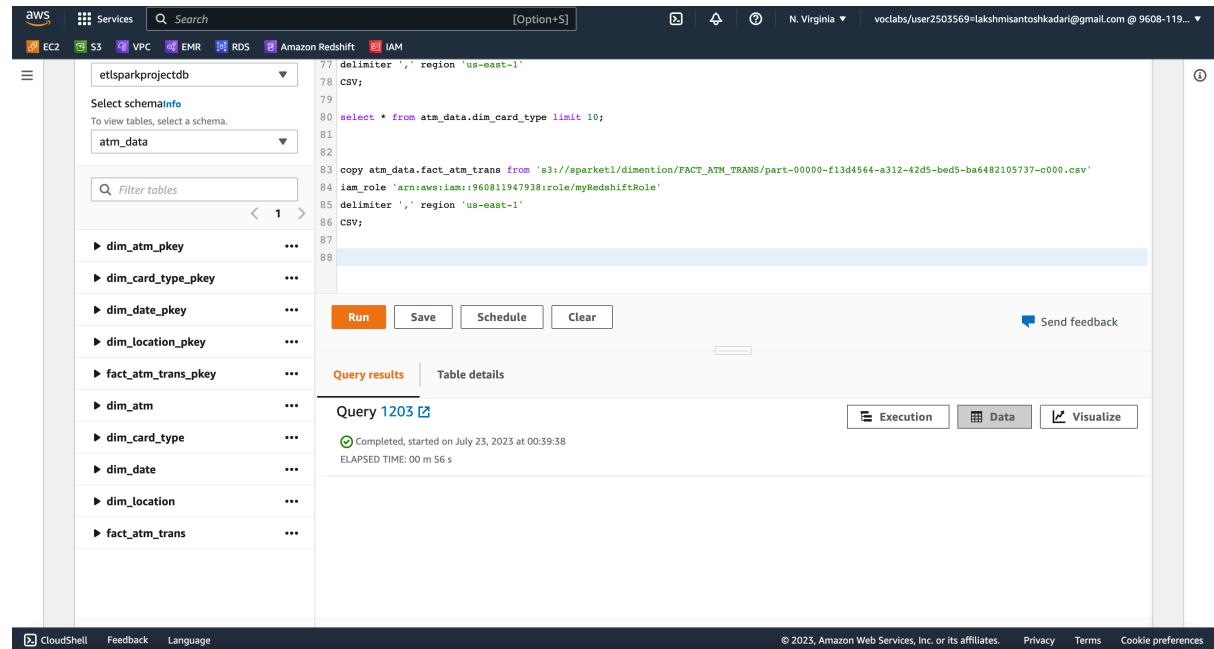
The screenshot shows the AWS Redshift Query Editor interface. The sidebar remains the same, showing 'etisparkprojectdb' for database and 'atm_data' for schema. The main area contains the following SQL code:

```
69 delimiter ',' region 'us-east-1'  
70 timeformat 'YYYY-MM-DDTHH:MI:SS'  
71 CSV;  
72  
73 select * from atm_data.dim_date limit 10;  
74  
75 copy atm_data.dim_card_type from 's3://sparketl/dimention/DIM_CARD_TYPE/part-00000-1aae...  
76 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
77 delimiter ',' region 'us-east-1'  
78 CSV;  
79  
80 select * from atm_data.dim_card_type limit 10;
```

Below the code, there are 'Run', 'Save', 'Schedule', and 'Clear' buttons. To the right, there are tabs for 'Query results' (selected), 'Table details', 'Execution', 'Data', and 'Visualize'. The status bar at the bottom indicates 'Completed, started on July 23, 2023 at 00:36:19' and 'ELAPSED TIME: 00 m 07 s'. A 'Rows returned (10)' message is displayed above the results table.

Copying the data to fact_atm_trans table

```
copy atm_data.fact_atm_trans from 's3://sparketl/dimention/FACT_ATM_TRANS/part-00000-f13d4564-a312-42d5-bed5-ba6482105737-c000.csv'  
iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
delimiter ',' region 'us-east-1'  
CSV;
```



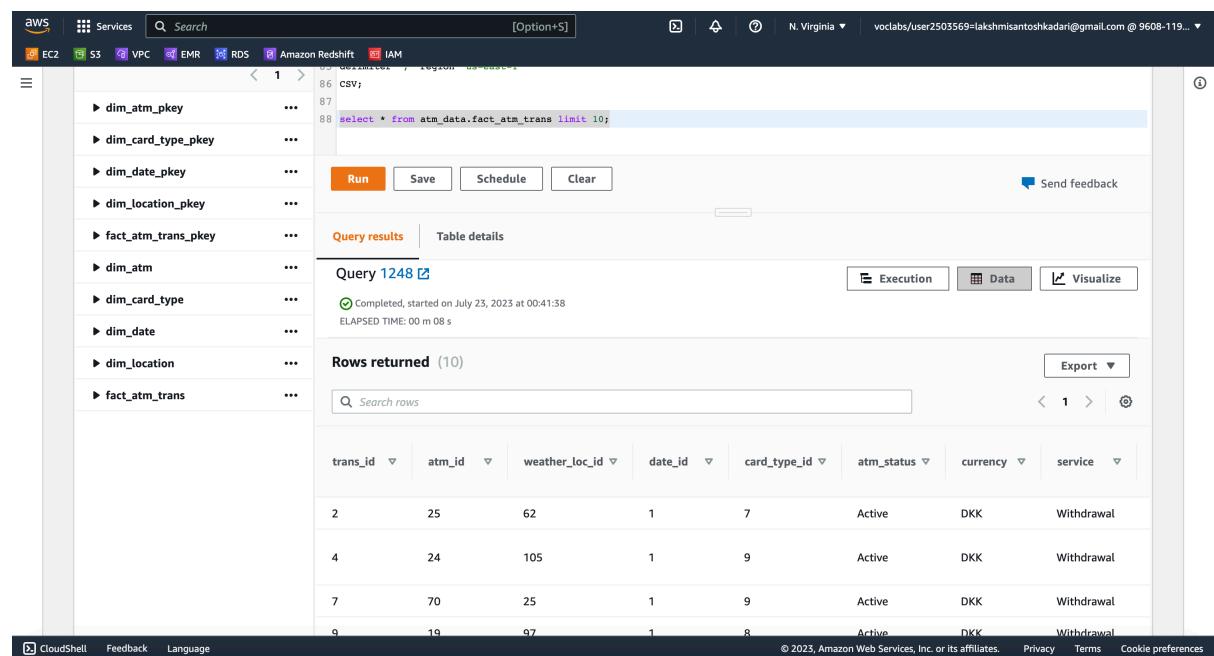
The screenshot shows the AWS Redshift console interface. In the left sidebar, there is a schema dropdown set to 'ettsparkprojectdb' and a table dropdown set to 'atm_data'. Below these dropdowns is a 'Filter tables' search bar. A list of tables is displayed, including 'dim_atm_pkey', 'dim_card_type_pkey', 'dim_date_pkey', 'dim_location_pkey', 'fact_atm_trans_pkey', 'dim_atm', 'dim_card_type', 'dim_date', 'dim_location', and 'fact_atm_trans'. The 'fact_atm_trans' table is currently selected. The main area contains a code editor with the following SQL query:

```
77 delimiter ',' region 'us-east-1'  
78 CSV;  
79  
80 select * from atm_data.dim_card_type limit 10;  
81  
82  
83 copy atm_data.fact_atm_trans from 's3://sparketl/dimention/FACT_ATM_TRANS/part-00000-f13d4564-a312-42d5-bed5-ba6482105737-c000.csv'  
84 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'  
85 delimiter ',' region 'us-east-1'  
86 CSV;  
87  
88
```

Below the code editor are four buttons: 'Run', 'Save', 'Schedule', and 'Clear'. To the right of the code editor is a 'Send feedback' link. Underneath the code editor, there is a 'Query results' section with a status message: 'Completed, started on July 23, 2023 at 00:39:38' and 'ELAPSED TIME: 00 m 56 s'. There are three tabs: 'Execution', 'Data', and 'Visualize', with 'Execution' being the active tab.

Verified the data to ensure data got uploaded using the below query.

```
select * from atm_data.fact_atm_trans limit 10;
```



The screenshot shows the AWS Redshift console interface. The left sidebar lists tables: 'dim_atm_pkey', 'dim_card_type_pkey', 'dim_date_pkey', 'dim_location_pkey', 'fact_atm_trans_pkey', 'dim_atm', 'dim_card_type', 'dim_date', 'dim_location', and 'fact_atm_trans'. The 'fact_atm_trans' table is selected. The main area contains a code editor with the following SQL query:

```
85 delimiter ',' region 'us-east-1'  
86 CSV;  
87  
88 select * from atm_data.fact_atm_trans limit 10;
```

Below the code editor are four buttons: 'Run', 'Save', 'Schedule', and 'Clear'. To the right of the code editor is a 'Send feedback' link. Underneath the code editor, there is a 'Query results' section with a status message: 'Completed, started on July 23, 2023 at 00:41:38' and 'ELAPSED TIME: 00 m 08 s'. There are three tabs: 'Execution', 'Data', and 'Visualize', with 'Execution' being the active tab. The 'Data' tab is active, showing a table titled 'Rows returned (10)'. The table has the following columns: trans_id, atm_id, weather_loc_id, date_id, card_type_id, atm_status, currency, and service. The data rows are:

trans_id	atm_id	weather_loc_id	date_id	card_type_id	atm_status	currency	service
2	25	62	1	7	Active	DKK	Withdrawal
4	24	105	1	9	Active	DKK	Withdrawal
7	70	25	1	9	Active	DKK	Withdrawal
9	19	97	1	8	Active	DKK	Withdrawal

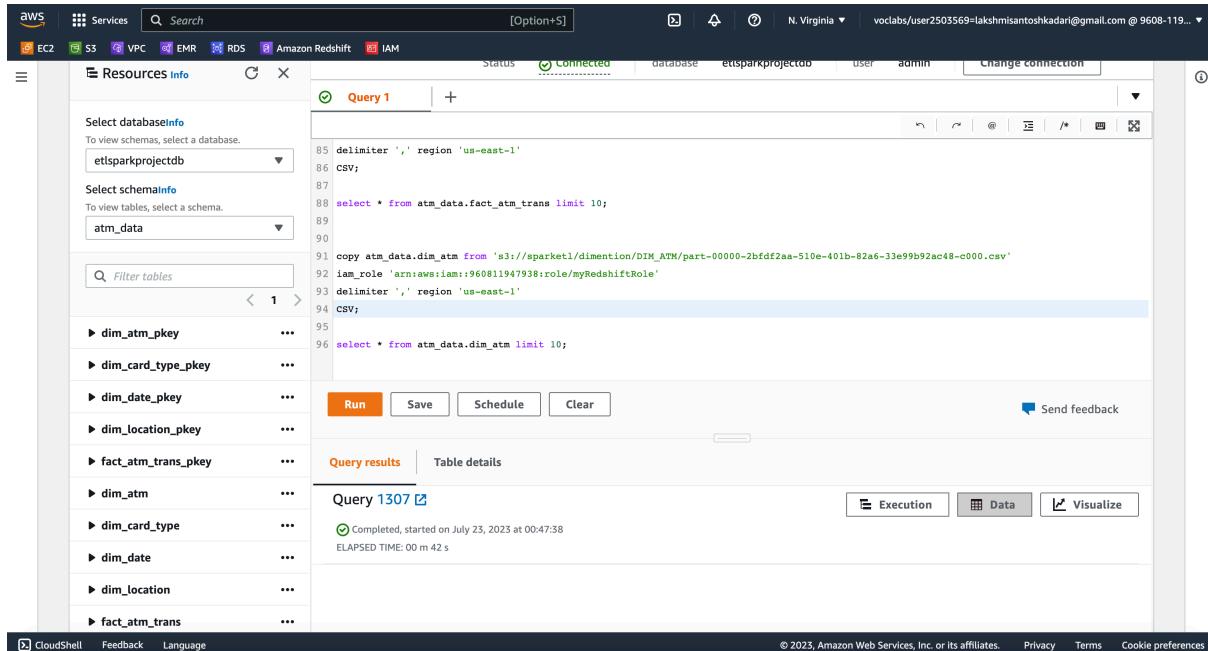
Copying the data to dim_atm table

```
copy atm_data.dim_atm from 's3://sparketl/dimention/DIM_ATM/part-00000-2bfdf2aa-510e-401b-82a6-33e99b92ac48-c000.csv'
```

```
iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
```

```
delimiter ',' region 'us-east-1'
```

CSV;



The screenshot shows the AWS Redshift Query Editor interface. The top navigation bar includes 'Services' (Search), 'Connected' status, database 'etlsparkprojectdb', user 'admin', and 'Change Connection'. The left sidebar lists tables: 'atm_data', 'dim_atm_pkey', 'dim_card_type_pkey', 'dim_date_pkey', 'dim_location_pkey', 'fact_atm_trans_pkey', 'dim_atm', 'dim_card_type', 'dim_date', 'dim_location', and 'fact_atm_trans'. The main area contains a query editor with the following SQL code:

```
85 delimiter ',' region 'us-east-1'
86 CSV;
87
88 select * from atm_data.fact_atm_trans limit 10;
89
90
91 copy atm_data.dim_atm from 's3://sparketl/dimention/DIM_ATM/part-00000-2bfdf2aa-510e-401b-82a6-33e99b92ac48-c000.csv'
92 iam_role 'arn:aws:iam::960811947938:role/myRedshiftRole'
93 delimiter ',' region 'us-east-1'
94 CSV;
```

Below the code, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted. To the right, there is a 'Send feedback' link. The results section shows 'Query 1307' completed at 00:47:38 with an elapsed time of 00 m 42 s. There are tabs for 'Execution', 'Data', and 'Visualize'.

Verified the data to ensure data got uploaded using the below query.

```
select * from atm_data.dim_atm limit 10;
```

aws Services Search [Option+S] N. Virginia v voclabs/user2503569=lakshmisantoshkadari@gmail.com @ 9608-119... ▾

EC2 S VPC EMR RDS Amazon Redshift IAM

☰ dim_atm_pkey ...
► dim_card_type_pkey ...
► dim_date_pkey ...
► dim_location_pkey ...
► fact_atm_trans_pkey ...
► dim_atm ...
► dim_card_type ...
► dim_date ...
► dim_location ...
► fact_atm_trans ...

95
96 select * from atm_data.dim_atm limit 10;

Run Save Schedule Clear Send feedback

Query results Table details

Query 1339 Execution Data Visualize

Completed, started on July 23, 2023 at 00:48:58
ELAPSED TIME: 00 m 08 s

Rows returned (10)

Export

Search rows < 1 >

atm_id	atm_number	atm_manufacturer	atm_location_id
5	102	NCR	3
15	111	Diebold Nixdorf	8
22	16	NCR	86
23	17	NCR	78
24	18	Diebold Nixdorf	105
27	20	NCR	14

CloudShell Feedback Language © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences